

Assignment 5: ML with sklearn

Name: Aarya Patil

Date: 04/05/23

```
In [ ]: # Import libraries
import pandas as pd
import seaborn as sb
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from sklearn.neural_network import MLPClassifier
```

1. Read the Auto data

```
In [ ]: # Read the data
df = pd.read_csv('Auto.csv')

# Output the first few rows
print(df.head(n = 5).to_string(index = False))

# Output the dimensions
print('\nDimensions:', df.shape)
```

mpg	cylinders	displacement	horsepower	weight	acceleration	year	orig
18.0	8	307.0	130	3504	12.0	70.0	
1	chevrolet	chevelle	malibu				
15.0	8	350.0	165	3693	11.5	70.0	
1	buick	skylark	320				
18.0	8	318.0	150	3436	11.0	70.0	
1	plymouth	satellite					
16.0	8	304.0	150	3433	12.0	70.0	
1	amc	rebel	sst				
17.0	8	302.0	140	3449	NaN	70.0	
1		ford	torino				

Dimensions: (392, 9)

2. Data exploration with code

```
In [ ]: # Use describe() on the mpg, weight, and year columns
print(df[["mpg", "weight", "year"]].describe(include = "all"))
```

	mpg	weight	year
count	392.000000	392.000000	390.000000
mean	23.445918	2977.584184	76.010256
std	7.805007	849.402560	3.668093
min	9.000000	1613.000000	70.000000
25%	17.000000	2225.250000	73.000000
50%	22.750000	2803.500000	76.000000
75%	29.000000	3614.750000	79.000000
max	46.600000	5140.000000	82.000000

Comments on the data exploration:

- Mpg -> Range: 37, Average: 23.445918
- Weight -> Range: 3527, Average: 2977.584184
- Year -> Range: 12, Average: 76.010256

3. Explore data types

```
In [ ]: # Check the data types of all columns
print('ORIGINAL DATATYPES:')
print(df.dtypes)

# Change the cylinders column to categorical (use cat.codes)
df['cylinders'] = df['cylinders'].astype('category').cat.codes

# Change the origin column to categorical (don't use cat.codes)
df['origin'] = df['origin'].astype('category')

# Verify the changes with the dtypes attribute
print('\nCHANGED DATATYPES:')
print(df.dtypes)
```

ORIGINAL DATATYPES:

```
mpg          float64
cylinders     int64
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        int64
name          object
dtype: object
```

CHANGED DATATYPES:

```
mpg          float64
cylinders     int8
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        category
name          object
dtype: object
```

4. Deal with NAs

```
In [ ]: # Delete rows with NAs
df = df.dropna()

# Output the new dimensions
print('New dimensions:', df.shape)
```

New dimensions: (389, 9)

5. Modify columns

```

In [ ]: # Make a new column, mpg_high (categorical)
high_or_low = [] # list to hold 0 or 1 values based on whether the mpg is h
i = 0 # counter variable

for item in df['mpg']: # assign whether the mpg is high(1) or low(0) by con
    if item > df['mpg'].mean():
        high_or_low.insert(i, 1)
    else:
        high_or_low.insert(i, 0)
    i = i+1

df['mpg_high'] = high_or_low # add the new column and assign the values to
df['mpg_high'] = df['mpg_high'].astype('category') # change the data type t
# df.dtypes # debug statement to make sure the type is changed

# Delete the mpg and name columns
df = df.drop(columns=['mpg', 'name'])

# Output the first few rows of the modified data frame
df.head()

```

```

Out[ ]:

```

	cylinders	displacement	horsepower	weight	acceleration	year	origin	mpg_high
0	4	307.0	130	3504	12.0	70.0	1	0
1	4	350.0	165	3693	11.5	70.0	1	0
2	4	318.0	150	3436	11.0	70.0	1	0
3	4	304.0	150	3433	12.0	70.0	1	0
6	4	454.0	220	4354	9.0	70.0	1	0

6. Data exploration with graphs

```


In [ ]: # Seaborn catplot on the mpg_high column
sb.catplot(data = df, x = 'mpg_high', kind = 'count')

```

```


Out[ ]: <seaborn.axisgrid.FacetGrid at 0x29d9de890>

```




```
In [ ]: # Seaborn relplot with horsepower on the x axis, weight on the y axis, and s
sb.relplot(data = df, x = 'horsepower', y = 'weight', hue = 'mpg_high')
```

```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x29d98ab00>
```



```
In [ ]: # Seaborn boxplot with mpg_high on the x axis and weight on the y axis  
sb.boxplot(data = df, x = 'mpg_high', y = 'weight')
```

```
Out[ ]: <Axes: xlabel='mpg_high', ylabel='weight'>
```



What I learned about the data from each graph:

- Catplot: I got to see how many cars have a high mpg vs. a low mpg.
- Relplot: I learned that as the weight of the car increases, so does the horsepower. The relationship they have is linear. Also cars that weigh less and have a lower horsepower tend to have a higher mpg.
- Boxplot: The second part that I learned from the relplot was confirmed in this graph. Cars that weigh less tend to have a higher mpg. I also got to see the range and average weight the cars with both high and low mpgs fall in.

7. Train/test split

```
In [ ]: # Train/test X data frames consists of all remaining columns except mpg_high
X = df.loc[:, ['cylinders', 'displacement', 'horsepower', 'weight', 'acceler
y = df['mpg_high']

# Split the training and testing data 80/20
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.8,

# Output the dimensions of train and test data
print('X train size:', X_train.shape)
print('X test size:', X_test.shape)
print('y train size:', y_train.shape)
print('y test size:', y_test.shape)

X train size: (311, 7)
X test size: (78, 7)
y train size: (311,)
y test size: (78,)
```

8. Logistic Regression

```
In [ ]: # Train a logistic regression model using solver lbfgs
logReg = LogisticRegression(solver = 'lbfgs', max_iter = 200)
logReg.fit(X_train, y_train)

# Test and evaluate
pred = logReg.predict(X_test)

# Print metrics using the classification report
print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.98	0.82	0.89	50
1	0.75	0.96	0.84	28
accuracy			0.87	78
macro avg	0.86	0.89	0.87	78
weighted avg	0.89	0.87	0.87	78

9. Decision Tree


```
In [ ]: # Train a decision tree
decTree = DecisionTreeClassifier()
decTree.fit(X_train, y_train)

# Test and evaluate
pred2 = decTree.predict(X_test)

# Print metrics using the classification report
print(classification_report(y_test, pred2))

# Plot the tree
print(tree.plot_tree(decTree))
```


	precision	recall	f1-score	support
0	0.92	0.90	0.91	50
1	0.83	0.86	0.84	28
accuracy			0.88	78
macro avg	0.87	0.88	0.88	78
weighted avg	0.89	0.88	0.89	78

```
[Text(0.6507352941176471, 0.9444444444444444, 'x[0] <= 2.5\ngini = 0.5\nsam
ples = 311\nvalue = [153, 158]'), Text(0.4338235294117647, 0.8333333333333333
34, 'x[2] <= 101.0\ngini = 0.239\nsamples = 173\nvalue = [24, 149]'), Text(
0.27941176470588236, 0.7222222222222222, 'x[5] <= 75.5\ngini = 0.179\nsampl
es = 161\nvalue = [16, 145]'), Text(0.14705882352941177, 0.6111111111111112
, 'x[1] <= 119.5\ngini = 0.362\nsamples = 59\nvalue = [14, 45]'), Text(0.05
8823529411764705, 0.5, 'x[4] <= 13.75\ngini = 0.159\nsamples = 46\nvalue =
[4, 42]'), Text(0.029411764705882353, 0.3888888888888889, 'gini = 0.0\nsamp
les = 2\nvalue = [2, 0]'), Text(0.08823529411764706, 0.3888888888888889, 'x
[3] <= 2683.0\ngini = 0.087\nsamples = 44\nvalue = [2, 42]'), Text(0.058823
529411764705, 0.2777777777777778, 'x[3] <= 2377.0\ngini = 0.045\nsamples =
43\nvalue = [1, 42]'), Text(0.029411764705882353, 0.16666666666666666, 'gin
i = 0.0\nsamples = 38\nvalue = [0, 38]'), Text(0.08823529411764706, 0.16666
666666666666, 'x[3] <= 2385.0\ngini = 0.32\nsamples = 5\nvalue = [1, 4]'),
Text(0.058823529411764705, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nv
alue = [1, 0]'), Text(0.11764705882352941, 0.05555555555555555, 'gini = 0.0
\nsamples = 4\nvalue = [0, 4]'), Text(0.11764705882352941, 0.27777777777777
78, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'), Text(0.23529411764705882, 0
.5, 'x[4] <= 17.75\ngini = 0.355\nsamples = 13\nvalue = [10, 3]'), Text(0.2
0588235294117646, 0.3888888888888889, 'x[2] <= 81.5\ngini = 0.469\nsamples
= 8\nvalue = [5, 3]'), Text(0.17647058823529413, 0.2777777777777778, 'gini
= 0.0\nsamples = 2\nvalue = [0, 2]'), Text(0.23529411764705882, 0.27777777
7777778, 'x[5] <= 71.5\ngini = 0.278\nsamples = 6\nvalue = [5, 1]'), Text(0
.20588235294117646, 0.16666666666666666, 'x[1] <= 131.0\ngini = 0.5\nsampl
es = 2\nvalue = [1, 1]'), Text(0.17647058823529413, 0.05555555555555555, 'gi
ni = 0.0\nsamples = 1\nvalue = [1, 0]'), Text(0.23529411764705882, 0.055555
55555555555, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'), Text(0.26470588235
29412, 0.16666666666666666, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'), Tex
```

```

t(0.2647058823529412, 0.3888888888888889, 'gini = 0.0\nsamples = 5\nvalue =
[5, 0]'), Text(0.4117647058823529, 0.6111111111111112, 'x[3] <= 3250.0\ngini
i = 0.038\nsamples = 102\nvalue = [2, 100]'), Text(0.35294117647058826, 0.5
, 'x[3] <= 2880.0\ngini = 0.02\nsamples = 100\nvalue = [1, 99]'), Text(0.32
35294117647059, 0.3888888888888889, 'gini = 0.0\nsamples = 94\nvalue = [0,
94]'), Text(0.38235294117647056, 0.3888888888888889, 'x[3] <= 2920.0\ngini
= 0.278\nsamples = 6\nvalue = [1, 5]'), Text(0.35294117647058826, 0.2777777
77777778, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'), Text(0.4117647058823
529, 0.277777777777778, 'gini = 0.0\nsamples = 5\nvalue = [0, 5]'), Text(0
.47058823529411764, 0.5, 'x[3] <= 3400.0\ngini = 0.5\nsamples = 2\nvalue =
[1, 1]'), Text(0.4411764705882353, 0.3888888888888889, 'gini = 0.0\nsamples
= 1\nvalue = [1, 0]'), Text(0.5, 0.3888888888888889, 'gini = 0.0\nsamples =
1\nvalue = [0, 1]'), Text(0.5882352941176471, 0.7222222222222222, 'x[4] <=
14.45\ngini = 0.444\nsamples = 12\nvalue = [8, 4]'), Text(0.558823529411764
7, 0.6111111111111112, 'x[5] <= 76.0\ngini = 0.444\nsamples = 6\nvalue = [2
, 4]'), Text(0.5294117647058824, 0.5, 'gini = 0.0\nsamples = 3\nvalue = [0,
3]'), Text(0.5882352941176471, 0.5, 'x[3] <= 2760.0\ngini = 0.444\nsamples
= 3\nvalue = [2, 1]'), Text(0.5588235294117647, 0.3888888888888889, 'gini =
0.0\nsamples = 2\nvalue = [2, 0]'), Text(0.6176470588235294, 0.38888888888
8889, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'), Text(0.6176470588235294,
0.6111111111111112, 'gini = 0.0\nsamples = 6\nvalue = [6, 0]'), Text(0.8676
470588235294, 0.8333333333333334, 'x[5] <= 79.5\ngini = 0.122\nsamples = 13
8\nvalue = [129, 9]'), Text(0.7941176470588235, 0.7222222222222222, 'x[4] <
= 21.6\ngini = 0.045\nsamples = 129\nvalue = [126, 3]'), Text(0.76470588235
29411, 0.6111111111111112, 'x[3] <= 2737.0\ngini = 0.031\nsamples = 128\nva
lue = [126, 2]'), Text(0.7058823529411765, 0.5, 'x[4] <= 13.45\ngini = 0.44
4\nsamples = 3\nvalue = [2, 1]'), Text(0.6764705882352942, 0.38888888888888
89, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'), Text(0.7352941176470589, 0.
3888888888888889, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'), Text(0.823529
4117647058, 0.5, 'x[2] <= 83.0\ngini = 0.016\nsamples = 125\nvalue = [124,
1]'), Text(0.7941176470588235, 0.3888888888888889, 'x[2] <= 79.5\ngini = 0.
375\nsamples = 4\nvalue = [3, 1]'), Text(0.7647058823529411, 0.27777777777
778, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'), Text(0.8235294117647058,
0.277777777777778, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'), Text(0.8529
411764705882, 0.3888888888888889, 'gini = 0.0\nsamples = 121\nvalue = [121,
0]'), Text(0.8235294117647058, 0.6111111111111112, 'gini = 0.0\nsamples = 1
\nvalue = [0, 1]'), Text(0.9411764705882353, 0.7222222222222222, 'x[6] <= 1
.5\ngini = 0.444\nsamples = 9\nvalue = [3, 6]'), Text(0.9117647058823529, 0
.6111111111111112, 'x[1] <= 247.0\ngini = 0.48\nsamples = 5\nvalue = [3, 2]
'), Text(0.8823529411764706, 0.5, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]')
), Text(0.9411764705882353, 0.5, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]')
, Text(0.9705882352941176, 0.6111111111111112, 'gini = 0.0\nsamples = 4\nva
lue = [0, 4]')]

```



10. Neural Network

```
In [ ]: # Train first neural network
neuralNet = MLPClassifier(hidden_layer_sizes = (7, 6, 5, 4, 3, 2, 1), random_state = 1234)
neuralNet.fit(X_train, y_train)

# Test and evaluate
pred3 = neuralNet.predict(X_test)

# Print metrics using the classification report
print(classification_report(y_test, pred3))

# Train second neural network
neuralNet2 = MLPClassifier(hidden_layer_sizes = (6, 3), random_state = 1234)
neuralNet2.fit(X_train, y_train)

# Test and evaluate
pred4 = neuralNet2.predict(X_test)

# Print metrics using the classification report
print(classification_report(y_test, pred4))
```

	precision	recall	f1-score	support
0	0.98	0.82	0.89	50
1	0.75	0.96	0.84	28
accuracy			0.87	78
macro avg	0.86	0.89	0.87	78
weighted avg	0.89	0.87	0.87	78

	precision	recall	f1-score	support
0	0.88	0.86	0.87	50
1	0.76	0.79	0.77	28
accuracy			0.83	78
macro avg	0.82	0.82	0.82	78
weighted avg	0.83	0.83	0.83	78

Model Comparison

- My first neural network model did much better than my second. I think that the overfitting in the first model somehow worked and gave me great precision. The accuracy isn't great for either model but it isn't terrible. Overall, the first model did better or equal in every category.

11. Analysis

My better neural network model and the logistic regression model both performed equally as well. The decision tree performed worse only in precision for 0 and recall for 0 but out-performed both of the other algorithms in all the other categories.

The accuracy for logistic regression and the neural network was 0.87 and for the decision tree it was 0.90. The recall for logistic regression and the neural network was 0.82(for 0) and 0.96(for 1) and for the decision tree it was 0.90. The precision for logistic regression and the neural network was 0.98(for 0) and 0.75(for 1) and for the decision tree it was 0.92(for 0) and 0.86(for 1).

I think that the data given to us was well suited for the decision tree. Usually I've noticed that logistic regression does really well, followed by neural network. I also think that maybe the network topology that I picked for the neural network models definitely skewed the data and caused it to return values that aren't as accurate. If I was to analyze this data again I would try to go about doing it differently so that I could get accurate results in all the algorithms.

I personally enjoy both R and sklearn pretty equally. I prefer R at times since many of the functions are built in but it's not too difficult in sklearn either. There are also things I prefer about sklearn though, such as the visual appearance, and the ease of using it. Overall, I would say they are fairly matched but if I had to choose only one to continue using, it would definitely be sklearn.